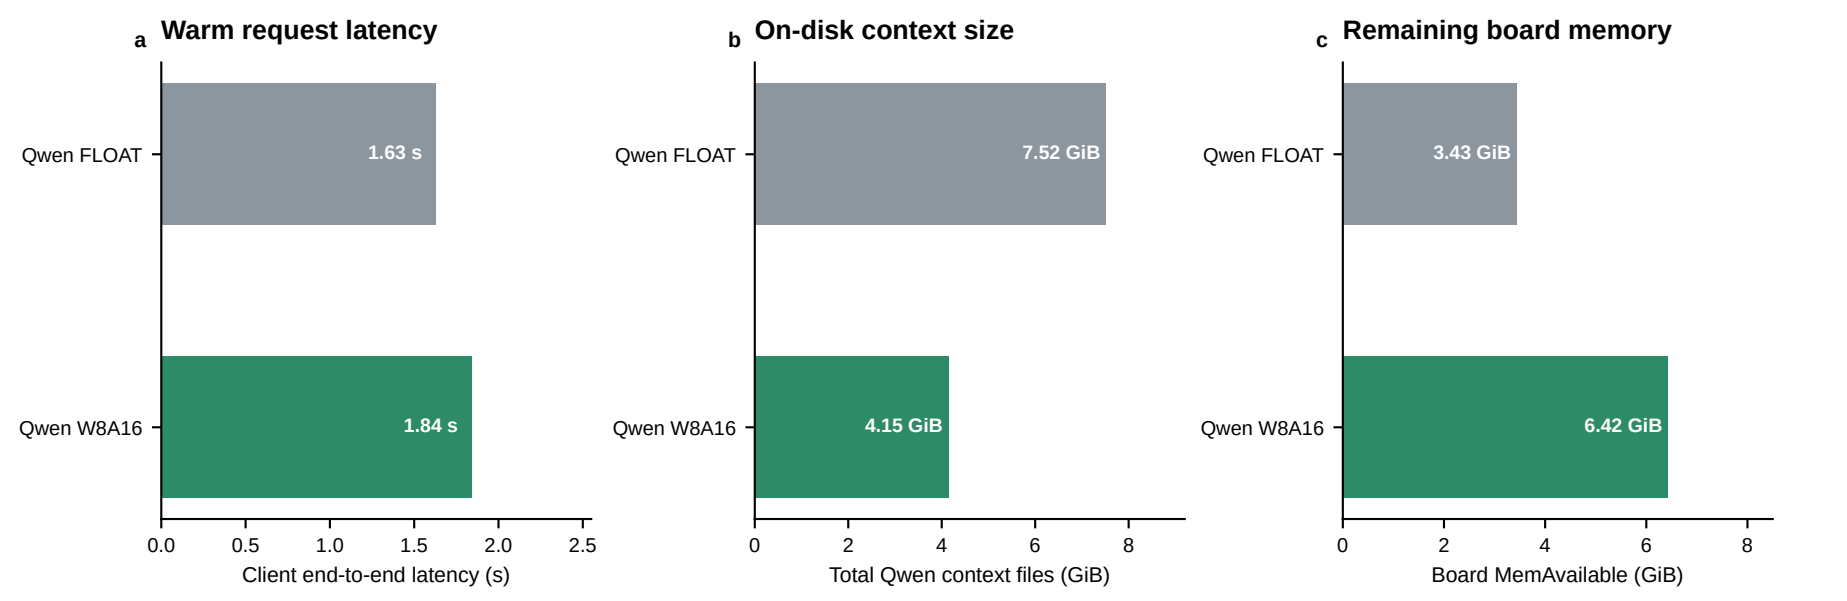


First cold request: 18.16 s → 7.70 s

W8A16 context files are 44.8% smaller.

W8A16 run: 2.99 GiB more available (warm median, n=23).

Paired QCS8550 test: 24 fixed requests; latency and board MemAvailable summaries exclude the first cold request (n=23).